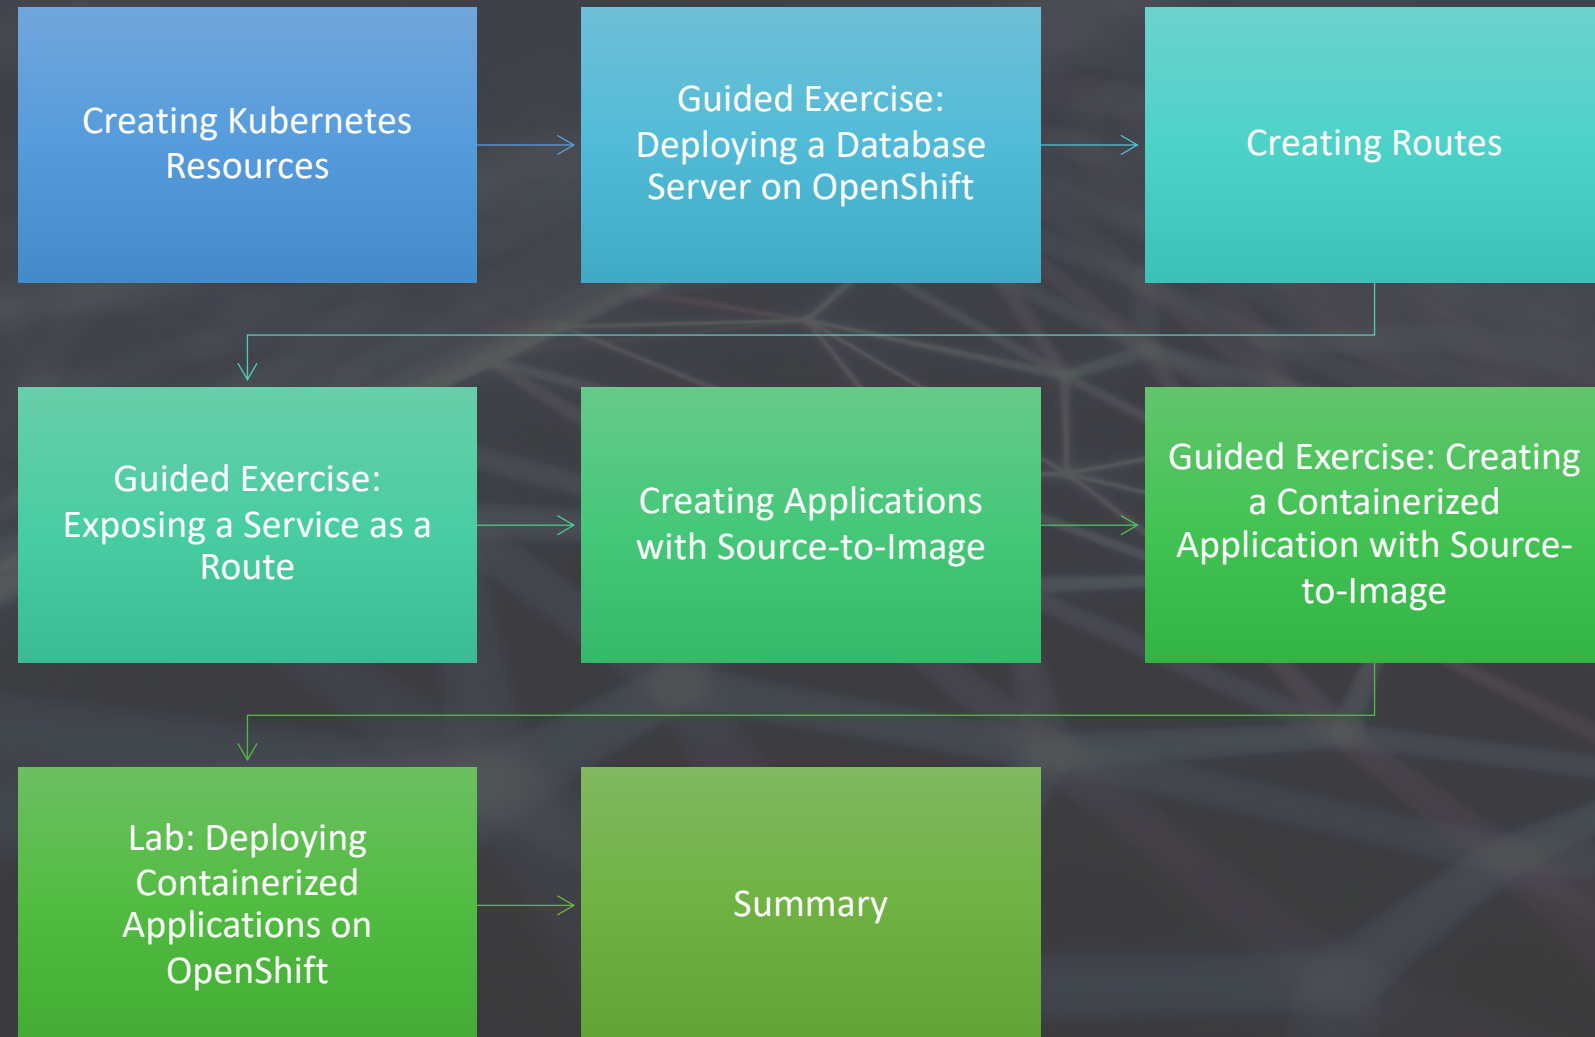


Describing K8s and OpenShift Architecture



Chapter objectives





Describing Kubernetes and OpenShift Architecture

After completing this section, you will be:

- Describe the architecture of a Kubernetes cluster running on the Red Hat OpenShift Container Platform (RHOCP)
- List the main resource types provided by Kubernetes and RHOCP
- Identify the network characteristics of containers, Kubernetes, and RHOCP
- List mechanisms to make a pod externally available

Kubernetes Terminology

Term	Definition
Node	A server that hosts applications in a Kubernetes cluster.
Master Node	A node server that manages the control plane in a Kubernetes cluster. Master nodes provide basic cluster services such as APIs or controllers.
Worker Node	Also named Compute Node , worker nodes execute workloads for the cluster. Application pods are scheduled onto worker nodes.
Resource	Resources are any kind of component definition managed by Kubernetes. Resources contain the configuration of the managed component (for example, the role assigned to a node), and the current state of the component (for example, if the node is available).
Controller	A controller is a Kubernetes process that watches resources and makes changes attempting to move the current state towards the desired state.
Label	A key-value pair that can be assigned to any Kubernetes resource. Selectors use labels to filter eligible resources for scheduling and other operations.
Namespace	A scope for Kubernetes resources and processes, so that resources with the same name can be used in different boundaries.



Note

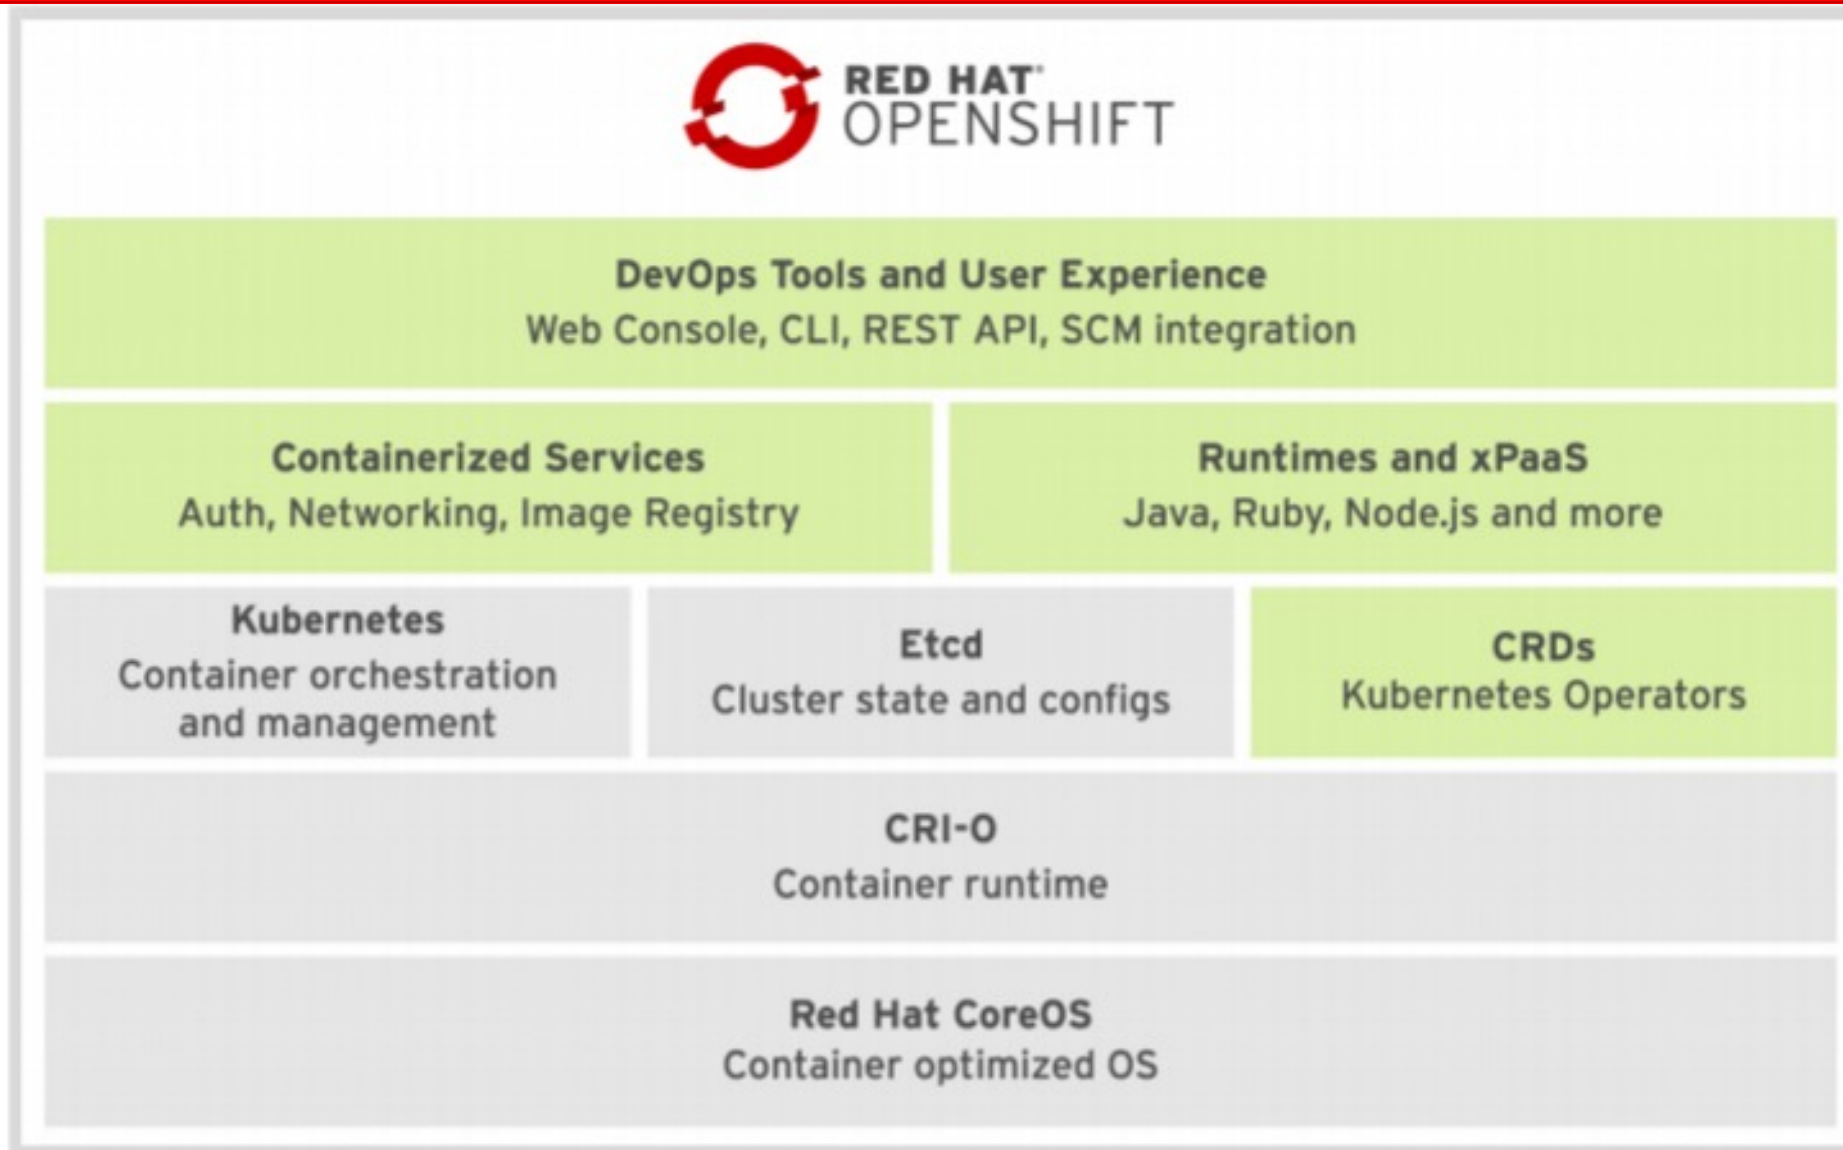
The latest Kubernetes versions implement many controllers as *Operators*. Operators are Kubernetes plug-in components that can react to cluster events and control the state of resources. Operators and CoreOS Operator Framework are outside the scope of this document.



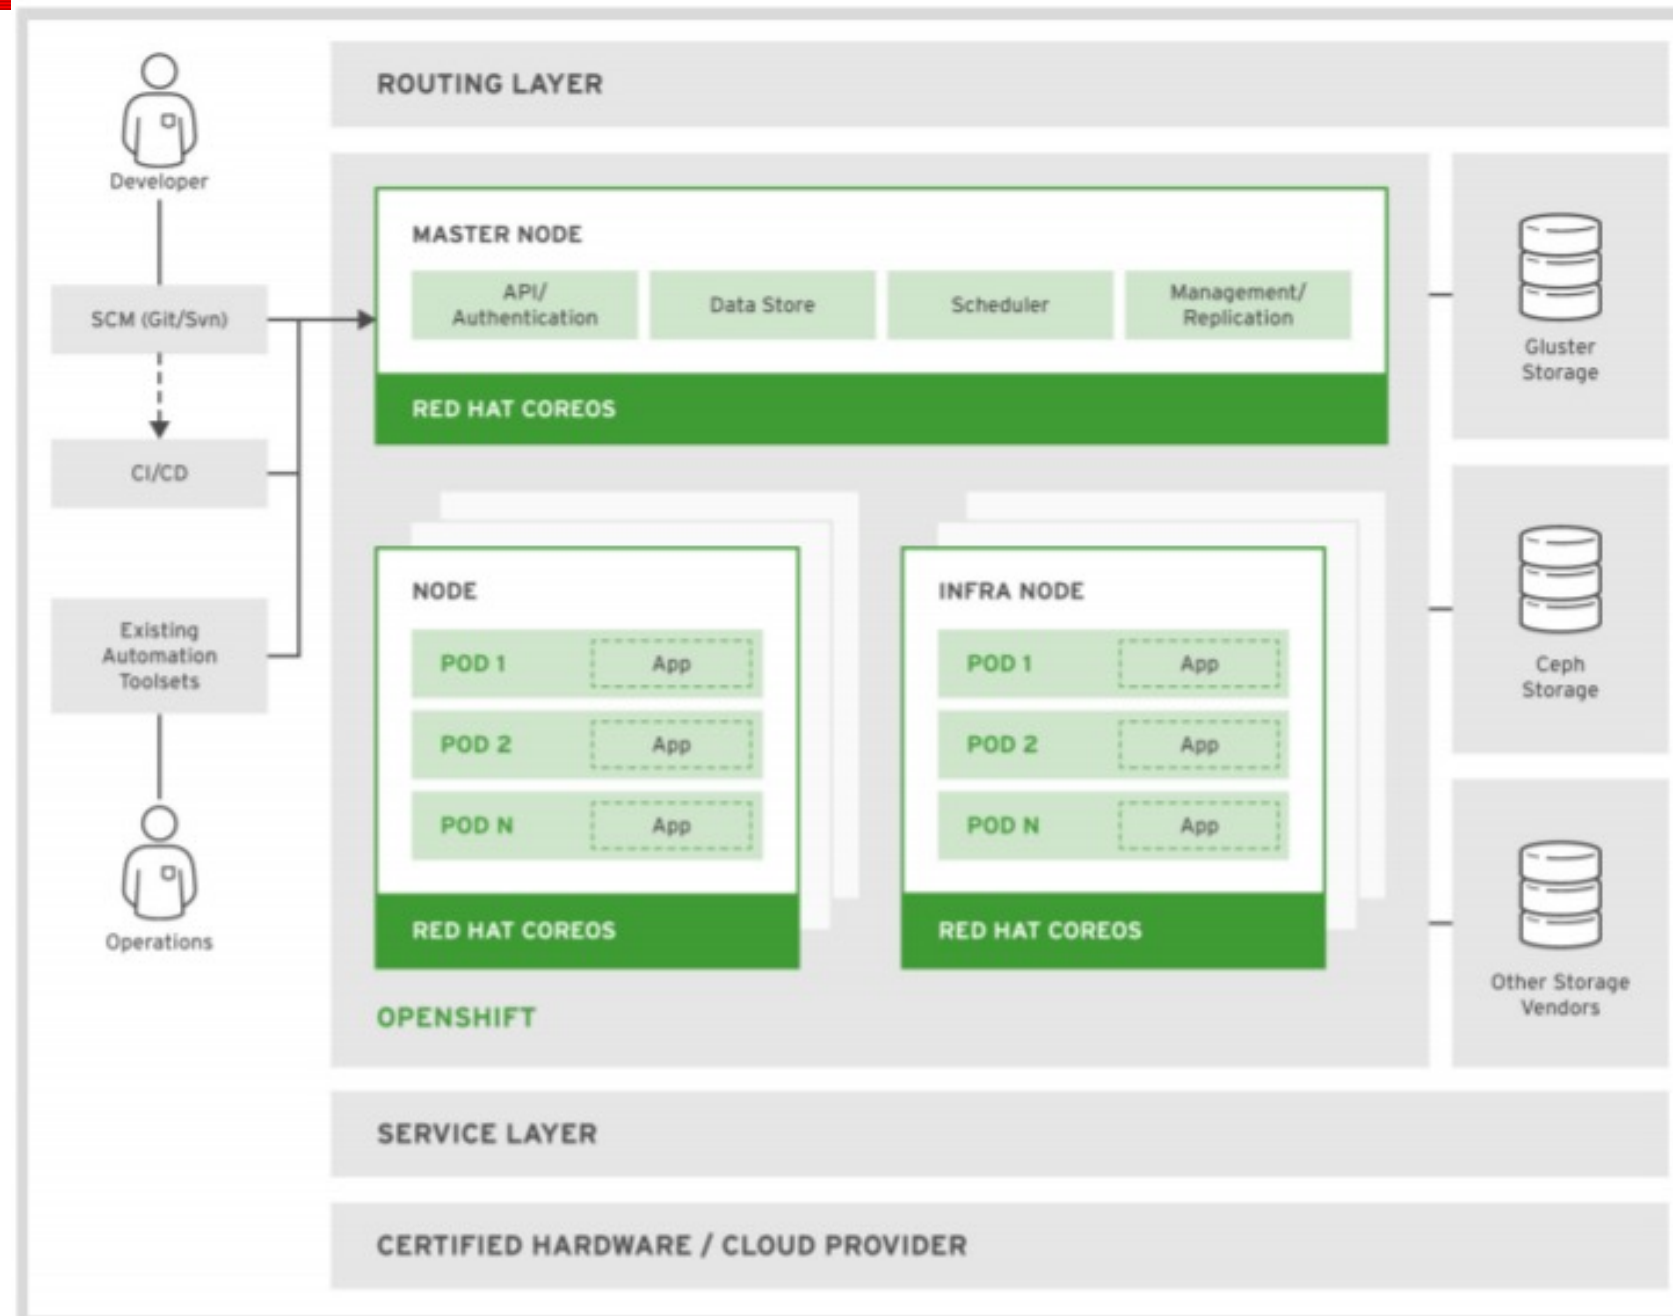
OpenShift Terminology

Term	Definition
Infra Node	A node server containing infrastructure services like monitoring, logging, or external routing.
Console	A web UI provided by the RHOCp cluster that allows developers and administrators to interact with cluster resources.
Project	OpenShift's extension of Kubernetes' namespaces. Allows the definition of user access control (UAC) to resources.

OpenShift Component Stack



OpenShift and K8s Architecture



New Features in RHOC v4

- CoreOS as mandatory operating system for all nodes
- Brand new cluster installer
- Self-managing platform
- Automatic cluster updates and recoveries
- Re-designed application life-cycle management
- Operator SDK: to build, test and package Operators

Describing K8s Resource Types

Pods (po)

Collection of containers that share resources. Basic unit of work for k8s

Services (svc)

Provide stable interface to all pods in deployment. By default, services connect clients to pods in round-robin fashion

Replication Controllers (rc)

A k8s resource that defines how pods are replicated and scheduled onto nodes. Basic k8s service to provide HA for pods and containers

Describing K8s Resource Types - continue

Persistent Volume (pv)

Provide permanent storage to pods

Persistent Volume Claims (pvc)

PVCs links PV to pod so that containers can mount storage to container's file system

ConfigMaps (cm) and Secrets

Contains set of keys values pair that can be used by resources.

Provides centralized configuration values used by resources.

Secrets values are always encoded and restricted to authorized users only

Describing OpenShift Resource Types

Deployment config(dc)

Like deployment, represents set of containers included in pod. Specifies deployment strategies to be used. A dc also provides basic but extensible continuous delivery workflow

Build Config (bc)

Defines process to be executed in project.

Used by S2I feature to build container image from application source code stored in Git repo.

A bc works with dc to provide basic but extensible continuous integration and continuous delivery workflow

Routes

Provides fully-qualified domain name recognized by OpenShift router as an ingress point for application and microservices

Deployment vs DeploymentConfig

```
apiVersion: v1
kind: DeploymentConfig
metadata:
  name: frontend
spec:
  replicas: 5
  selector:
    name: frontend
  template: { ... }
  triggers:
    - type: ConfigChange ①
    - imageChangeParams:
        automatic: true
        containerNames:
          - helloworld
        from:
          kind: ImageStreamTag
          name: hello-openshift:latest
        type: ImageChange ②
  strategy:
    type: Rolling ③
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello-openshift
spec:
  replicas: 1
  selector:
    matchLabels:
      app: hello-openshift
  template:
    metadata:
      labels:
        app: hello-openshift
    spec:
      containers:
        - name: hello-openshift
          image: openshift/hello-openshift:latest
          ports:
            - containerPort: 80
```

DC add enhancement to deployment.

- **Replication Controllers:** such as ability to involve one or more replication controllers, using pod template in controlling point-in-time record state of pods
- **Triggers:** add on triggers that drive automated deployments in response to events (Image change or Config change)
- **Custom strategies:** user-customizable deployment strategies to transition from previous version to new version.
- **Lifecycle hooks:** Uses hooks or lifecycle hooks for executing custom behaviour after pods are being built
- **Automatic rollbacks :** versioning of app in order to support rollbacks either automatically or manually in case of failure during deployment

Networking

Container IP

Ephemeral IP address

Assigned from internal network that is accessible only from node running the container

Software-Defined network (SDN)

Provides communication between container in pods between nodes

Access to SDN only works from inside same Kubernetes cluster

Networking - continue

Services

Containers do not connect each other dynamic IP address directly

Uses services by linking more stable IP addresses from SDN to pods

Pods restarted, replicated, rescheduled to different nodes – services get updated, providing scalability and high availability

External Access / K8s Ingress

Is more complicated.

K8s uses NodePort attribute to provide external access. But it is insecure and doesn't scale well

Networking - continue

OpenShift Routes - External Access

Simpler by defining route resources

Route defines external-facing DNS names and ports for service

A router (ingress controller) forwards HTTP(s) requests to service addresses inside K8s SDN.

OpenShift map IP addresses of RHOCN router nodes

Features	K8s Ingress	OpenShift Route
Standard Kubernetes object	X	
External access to services	X	X
Persistent (sticky) sessions	X	X
Load-balancing strategies (e.g. round robin)	X	X
Rate-limit and throttling	X	X
IP whitelisting	X	X
TLS edge termination for improved security	X	X
TLS re-encryption for improved security		X
TLS passthrough for improved security		X
Multiple weighted backends (split traffic)		X
Generated pattern-based hostnames		X
Wildcard domains		X

Quiz 1

Which two sentences are correct regarding Kubernetes architecture? (Choose two)

- a) Kubernetes nodes can be managed without a master
- b) Kubernetes masters schedule pods to specific nodes
- c) Kubernetes tools cannot be used to manage resources in an OpenShift cluster
- d) Kubernetes masters manage pod scaling
- e) Containers created from Kubernetes pods cannot be managed using standalone tools such as Podman

Quiz 1

Which two sentences are correct regarding Kubernetes architecture? (Choose two)

- a) Kubernetes nodes can be managed without a master
- b) Kubernetes masters schedule pods to specific nodes
- c) Kubernetes tools cannot be used to manage resources in an OpenShift cluster
- d) Kubernetes masters manage pod scaling
- e) Containers created from Kubernetes pods cannot be managed using standalone tools such as Podman

Quiz 2

Which two sentences are correct regarding Kubernetes and OpenShift resource types? (Choose two)

- a) A pod is responsible for provisioning its own persistent storage
- b) All pods generated from the same replication controller have to run in the same node
- c) A route is responsible for providing IP addresses for external access to pods
- d) A replication controller is responsible for monitoring and maintaining the number of pods for a particular application

Quiz 2

Which two sentences are correct regarding Kubernetes and OpenShift resource types? (Choose two)

- a) A pod is responsible for provisioning its own persistent storage
- b) All pods generated from the same replication controller have to run in the same node
- c) A route is responsible for providing IP addresses for external access to pods
- d) A replication controller is responsible for monitoring and maintaining the number of pods for a particular application

Quiz 3

Which two statements are true regarding Kubernetes and OpenShift networking?

(Choose two)

- a) A Kubernetes service can provide an IP address to access a set of pods
- b) Kubernetes is responsible for providing a fully qualified domain name for a pod
- c) A replication controller is responsible for routing external requests to the pods
- d) A route is responsible for providing DNS names for external access

Quiz 3

Which two statements are true regarding Kubernetes and OpenShift networking?
(Choose two)

- a) A Kubernetes service can provide an IP address to access a set of pods
- b) Kubernetes is responsible for providing a fully qualified domain name for a pod
- c) A replication controller is responsible for routing external requests to the pods
- d) A route is responsible for providing DNS names for external access

Quiz 4

Which statement is correct regarding persistent storage in Kubernetes and OpenShift?

- a) A PVC represents a storage area that a pod can use to store data and is provisioned by the application developer
- b) A PVC represents a storage area that can be requested by a pod to store data but is provisioned by the cluster administrator
- c) A PVC represents the amount of memory that can be allocated to a node, so that a developer can state how much memory he requires for his application to run
- d) A PVC represents the number of CPU processing units that can be allocated to pod

Quiz 4

Which statement is correct regarding persistent storage in Kubernetes and OpenShift?

- a) A PVC represents a storage area that a pod can use to store data and is provisioned by the application developer
- b) A PVC represents a storage area that can be requested by a pod to store data but is provisioned by the cluster administrator
- c) A PVC represents the amount of memory that can be allocated to a node, so that a developer can state how much memory he requires for his application to run
- d) A PVC represents the number of CPU processing units that can be allocated to pod

Quiz 5

Which statement is correct regarding OpenShift additions to Kubernetes?

- a) OpenShift adds features to simplify Kubernetes configuration of many real-world use cases
- b) Container images created for OpenShift cannot be used with plain Kubernetes
- c) Red Hat maintains forked versions of Kubernetes internal to the RHOCP product
- d) Doing continuous integration and continuous deployment with RHOCP requires external tools

Quiz 5

Which statement is correct regarding OpenShift additions to Kubernetes?

- a) OpenShift adds features to simplify Kubernetes configuration of many real-world use cases
- b) Container images created for OpenShift cannot be used with plain Kubernetes
- c) Red Hat maintains forked versions of Kubernetes internal to the RHOCp product
- d) Doing continuous integration and continuous deployment with RHOCp requires external tools



Creating Kubernetes Resources

After completing this section, you will be:

- Able to create standard K8s resources

The OpenShift Command

- Main method to interact with RHOCF cluster

```
$ oc <command> --parameters ...
```

- Requires logged-in user to cluster

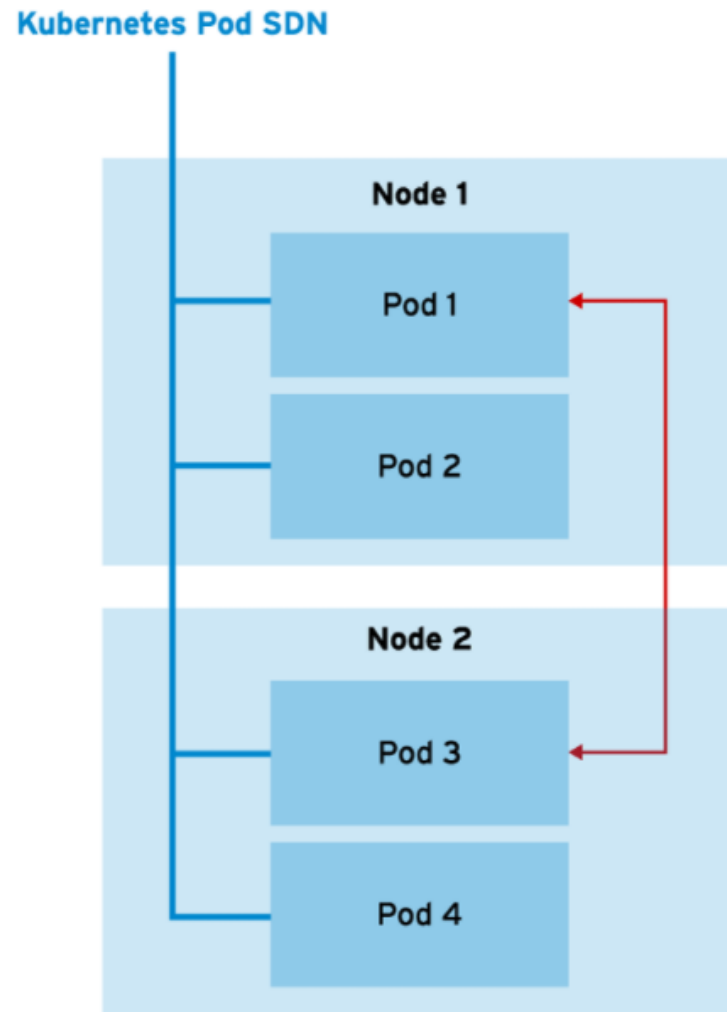
```
$ oc login -u <username> -p <password> <cluster-api-Url>
```

Describing Pod Resource Definition Syntax

```
apiVersion: v1
kind: Pod1
metadata:
  name: wildfly2
  labels:
    name: wildfly3
spec:
  containers:
    - resources:
        limits :
          cpu: 0.5
      image: do276/todojee
      name: wildfly
      ports:
        - containerPort: 80804
          name: wildfly
      env:5
        - name: MYSQL_ENV_MYSQL_DATABASE
          value: items
        - name: MYSQL_ENV_MYSQL_USER
          value: user1
        - name: MYSQL_ENV_MYSQL_PASSWORD
          value: mypa55
```


1. Declares a Kubernetes pod resource type.
2. A unique name for a pod in Kubernetes that allows administrators to run commands on it.
3. Creates a label with a key value pair
4. A container-dependent attribute identifying which port on the container is exposed.
5. Defines a collection of environment variables.

SDN exists on each node.



- Each node hosts own virtual network
- OS Kernel provides IP address to Nodes
- CNI provides IP address to Pods
- Pods from same node able to communicate
- Pods from different node can't communicate
- Nodes' subnet will always differs from Pods' subnet

Services resource type



- Provides a stable source IP address
- Each pod matching selector is added to services as endpoint
- Can load balance incoming traffic
- Expose service to get FQDN route to Pods

Describing Service Resource Definition Syntax

```
{
  "kind": "Service",
  "apiVersion": "v1",
  "metadata": {
    "name": "quotedb"
  },
  "spec": {
    "ports": [
      {
        "port": 3306,
        "targetPort": 3306
      }
    ],
    "selector": {
      "name": "mysqldb"
    }
  }
}
```

1. The kind of Kubernetes resource. In this case, a Service.
2. A unique name for the service.
3. ports is an array of objects that describes network ports exposed by the service.
4. service port and the service forwards packets to the pod targetPort.
5. selector is how the service finds pods to forward packets to.

Discovering Services

- Pods find a service IP address and port using environment variables.
- OpenShift automatically injects into containers for all pods inside same project
 - `SVC_NAME_SERVICE_HOST` is the service IP address.
 - `SVC_NAME_SERVICE_PORT` is the service TCP port.
- OpenShift internal DNS server
 - `SVC_NAME.PROJECT_NAME.svc.cluster.local`



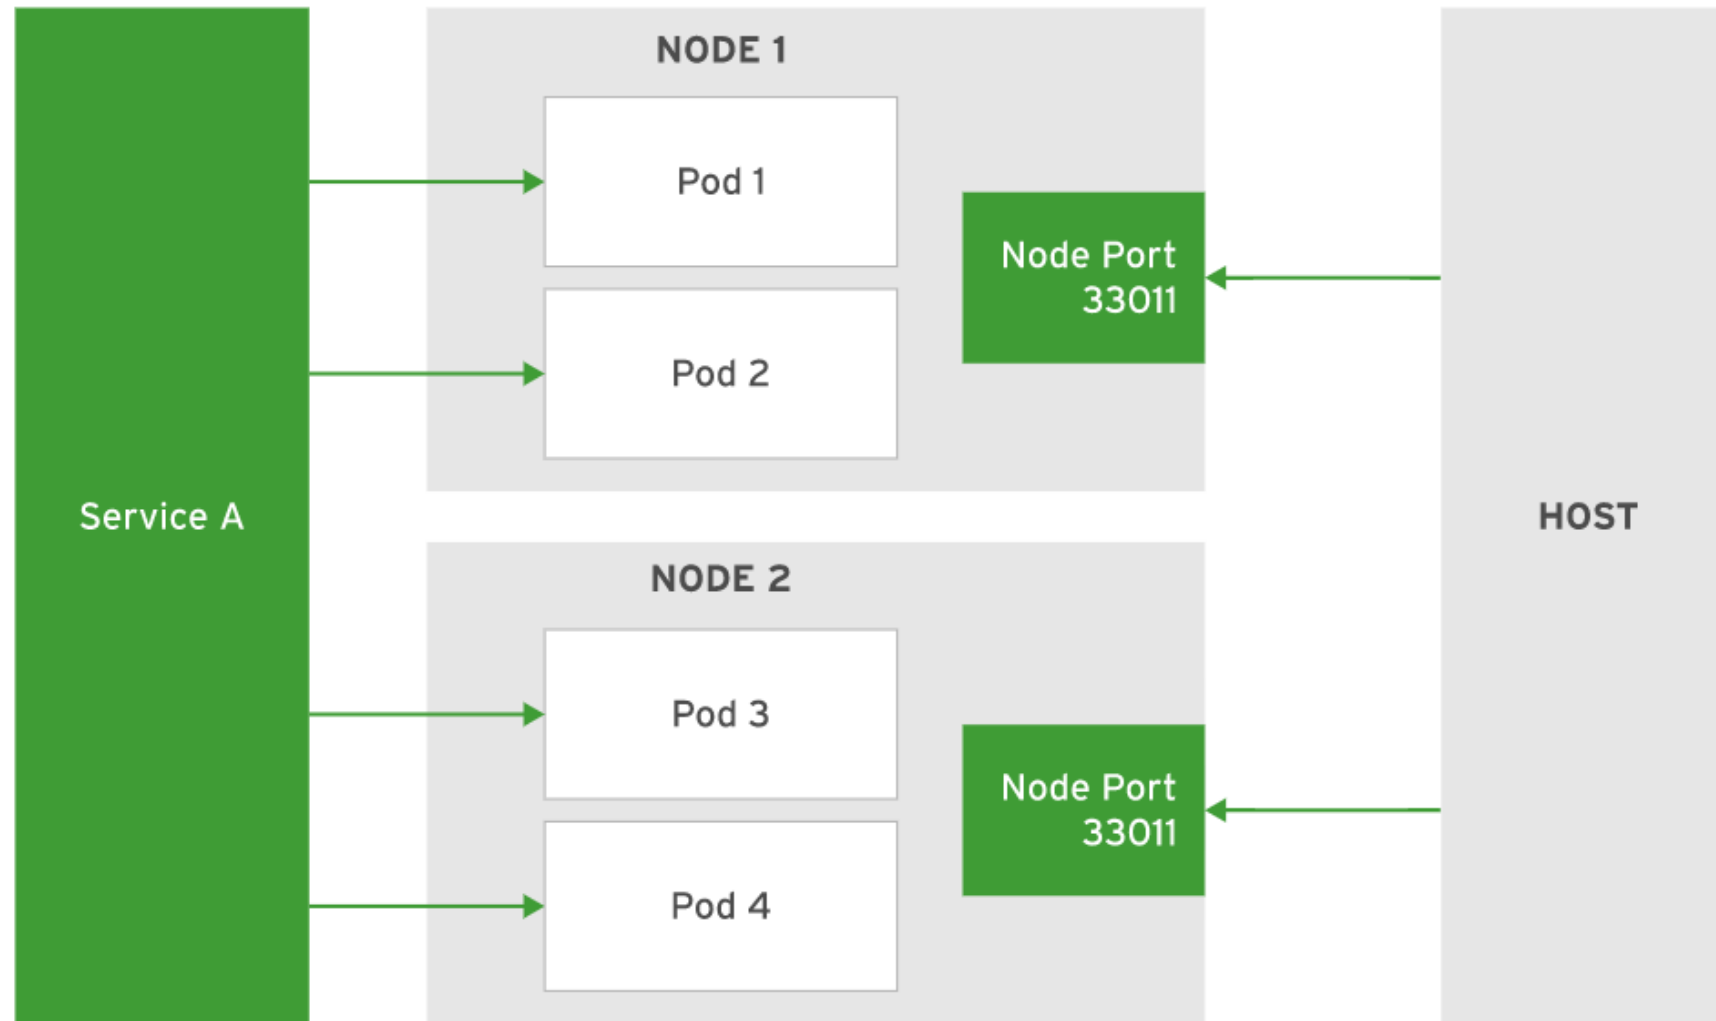
Note

The `SVC_NAME` part of the variable is changed to comply with DNS naming restrictions: letters are capitalized and underscores (`_`) are replaced by dashes (`-`).

Providing access to service from outside

- NodePort type:
 - Bind worker node port to service port
 - Use `oc edit svc` subcommand and specify NodePort as type
 - Insecured and inefficient
 - Network Port : 30000 - 37000
- OpenShift Routes:
 - Preferred approach
 - Expose services create unique URL (route)
 - Use `oc expose` subcommand or OpenShift web UI

Use NodePort type:



Use `oc port-forward` subcommand

- Mapping forwards connections to single pod.
- Different from having access via service resource
- Lack of many features offered by service resource
- Mapping exists only on workstation where `oc` is executed.
- For testing, troubleshooting, diagnose networking access

```
[student@workstation ~] $ oc port-forward mysql-openshift-1-glqrp 3306:3306
```

Compares to `oc expose services` subcommand

- Mappings exists for all network users
- Load-balances connections to potentially multiple pods

Creating New Applications

- Use `oc create` subcommand
 - Create resources by definition file or stdin JSON/YAML
- Use `oc apply` subcommand
 - Merge existing resources definition with stdin JSON/YAML
 - Create resources if it doesn't exist
- Use `oc new-app` command
 - Native to OpenShift
 - More dynamic
 - Able to use different method:
 - docker-image
 - Source-code (local or remote Git repository)
 - Template

Use `oc new-app` subcommand - Example

- Create an application based on a image from quay.io:

```
$ oc new-app --name app1 --docker-image quay.io/jason.wong76/webserver  
$ oc status
```

- Create an application using source code from github:

```
$ oc new-app --name app2 ruby~https://github.com/sclorg/ruby-ex.git --as-deployment-config  
$ oc status
```

- Use openshift-template:

```
$ oc new-app --name app3 --template=mysql-ephemeral -n openshift  
$ oc status
```

Use oc new-app subcommand - Example

- Create an application based on a mysql from Docker Hub with label and application's parameters

```
$ oc new-app --name mydb --as-deployment-config \  
    -e MYSQL_USER=redhat -e MYSQL_PASSWORD=redhat \  
    -e MYSQL_ROOT_PASSWORD=redhat -e MYSQL_DATABASE=datadb \  
    -l db=hrapp1 mysql
```

-e application_variable = data

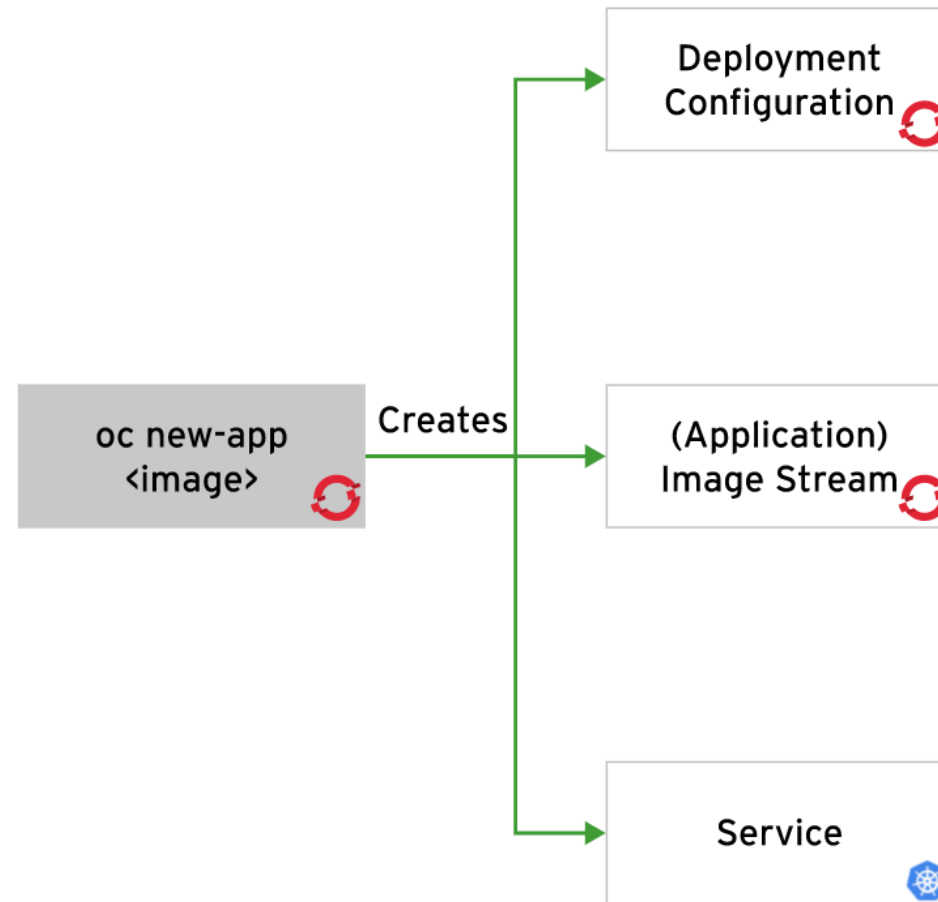
-l label = data

mysql : image stream

--name application name

--as-deployment-config : native to OCP

Resources created by `oc new-app` subcommand



→ created by `oc new-app`

Managing OpenShift Resources at the Command Line

- Use `oc get` subcommand to retrieve information about resources

- The syntax:

```
$ oc get RESOURCE_TYPE [-o wide] [-o yaml [ > output_file.yaml ] ]
```

- Resource_type:

- pod
- deployment
- deploymentconfig or dc
- service or svc
- route
- build config or bc
- imagestream or is

Managing OpenShift Resources at the Command Line

- Use `oc get all` subcommand to retrieve summary info on all resources

```
$ oc get all
```

NAME	DOCKER REPO	TAGS	UPDATED
is/nginx	172.30.1.1:5000/basic-kubernetes/nginx	latest	About an hour ago

NAME	REVISION	DESIRED	CURRENT	TRIGGERED BY
dc/nginx	1	1	1	config,image/nginx:latest

NAME	DESIRED	CURRENT	READY	AGE
rc/nginx-1	1	1	1	1h

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
svc/nginx	172.30.72.75	<none>	80/TCP,443/TCP	1h

NAME	READY	STATUS	RESTARTS	AGE
po/nginx-1-ypp8t	1/1	Running	0	1h

Managing OpenShift Resources at the Command Line

- Use `oc describe` subcommand to retrieve details information about resources

- The syntax:

```
$ oc describe RESOURCE_TYPE
```

- Resource_type:
 - pod
 - deployment
 - deployment config or dc
 - service or svc
 - route
 - build config or bc
 - imagestream or is

Managing OpenShift Resources at the Command Line

- Retrieve detailed information on particular pod

```
$ oc describe pod mysql-openshift-1-glqrp
```

```
Name:          mysql-openshift-1-glqrp
Namespace:     mysql-openshift
Priority:       0
PriorityClassName: none
Node:          cluster-worker-1/172.25.250.52
Start Time:    Fri, 15 Feb 2019 02:14:34 +0000
Labels:        app=mysql-openshift
               deployment=mysql-openshift-1
               deploymentconfig=mysql-openshift
Annotations:   openshift.io/deployment-config.latest-version: 1
               openshift.io/deployment-config.name: mysql-openshift
               openshift.io/deployment.name: mysql-openshift-1
               openshift.io/generated-by: OpenShiftNewApp
               openshift.io/scc: restricted
Status:        Running
IP:            10.129.0.85
```

Managing OpenShift Resources at the Command Line

- Use `oc edit` subcommand to edit resource definition. uses vi

```
$ oc edit dc/mysql
```

- Or use `oc patch` subcommand to edit resource definition. Direct.

```
$ oc patch node node1 -p '{"spec":{"unschedulable":true}}'
```

- Use `oc delete` subcommand to delete resource.

```
$ oc delete RESOURCE_TYPE <resource_name>
```

- Execute commands inside pod named `app1-n9h1t`

```
$ oc exec app1-n9h1t /bin/bash  
$ oc exec app1-n9h1t cat /etc/*release
```

- Delete a pod named `app1-n9h1t`

```
$ oc delete pod app1-n9h1t
```

Labelling resources

- Grouping resources by application, environment or some other criteria.
- Labels are part of metadata section of resource
- Defined as key/value pairs

```
apiVersion: v1
kind: Service
metadata:
  ...contents omitted...
  labels:
    app: nexus
    template: nexus-persistent-template
  name: nexus
  ...contents omitted...
```

Use -l option to use label

- Retrieve information on all resources belong to label `app=nexus`

```
$ oc get svc,dc -l app=nexus
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/nexus	ClusterIP	172.30.29.218	<none>	8081/TCP	4h

NAME	REVISION	DESIRED	CURRENT	...
deploymentconfig.apps.openshift.io/nexus	1	1	1	...

- Both `app` and `template` key labels are common
- The `app` key indicates application
- The `template` key labels all resources generated by the same template name

Explore it further!

- Creating Routes
- Creating Applications with Source-to-Image
- Creating Applications with OpenShift Web Console

Chapter Summary

In this chapter, you learned:




Use the OpenShift command-line client `oc`



Source-to-Image (S2I)



The `oc new-app` command can create application pods



A Route connects a public-facing IP address and DNS host name to an internal-facing service IP